



Robust Neural Learning via S-Divergence — Complete Honest Benchmark (seed=42, single run)
 ★ SDIV(opt) $\beta=0.05, \lambda=-0.4$ = best clean acc; key finding: λ controls convergence regime

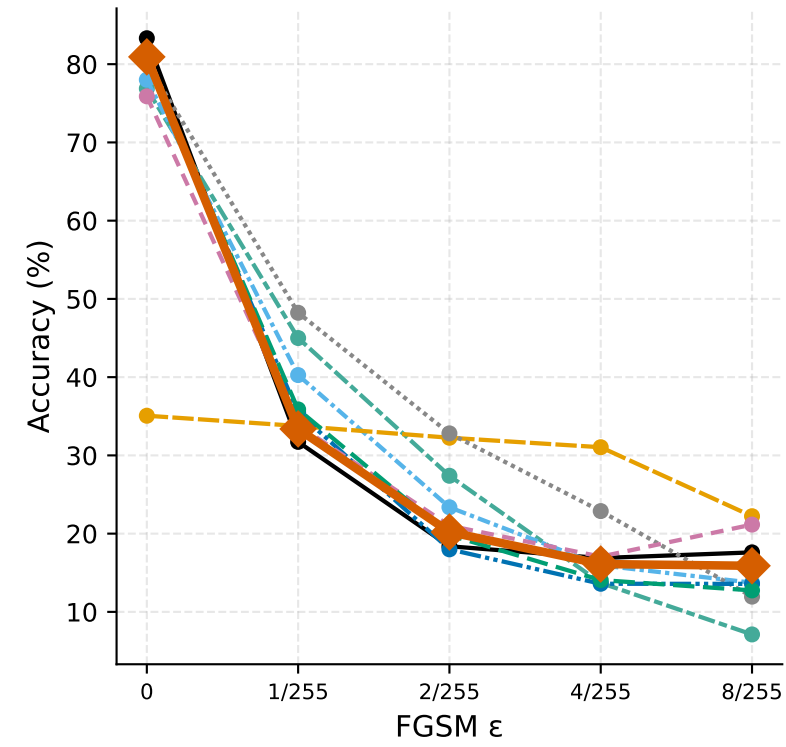
(a) PathMNIST Noise Robustness



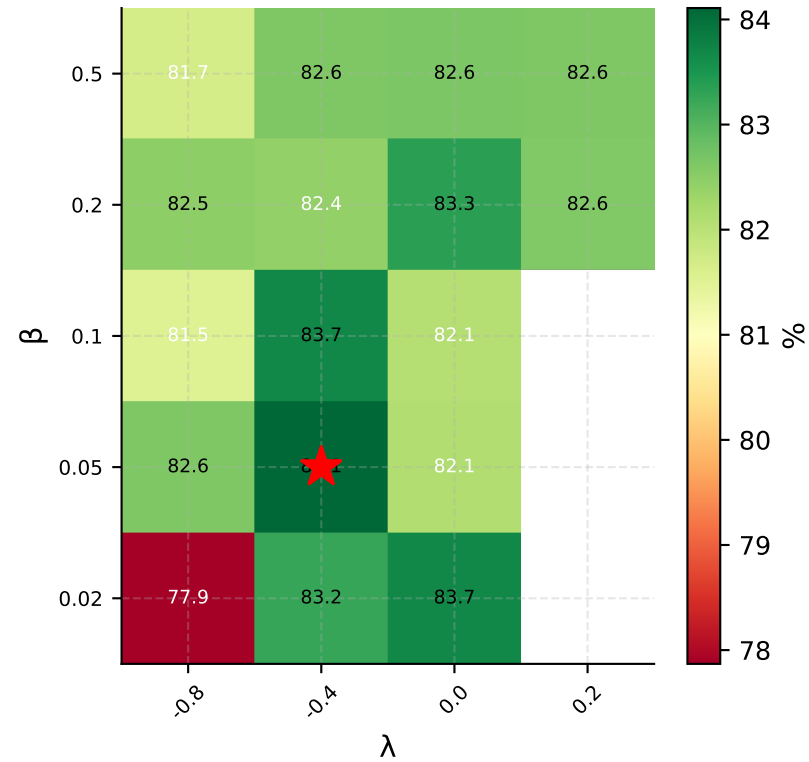
(b) DermaMNIST Noise
 ($\triangle$ trivial solutions present)



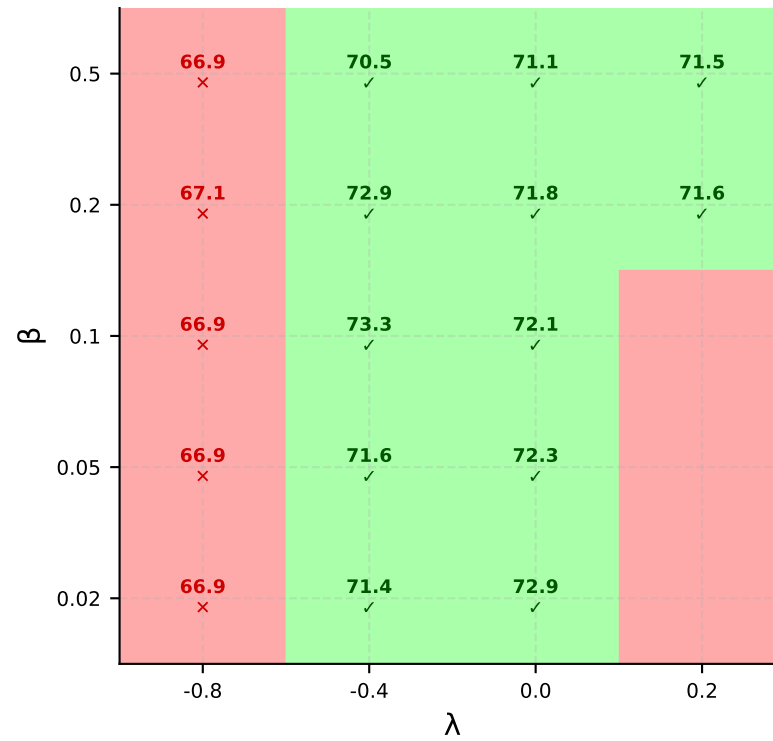
(c) PathMNIST FGSM
 (all losses vulnerable)



(d) PathMNIST (β, λ) Grid
 ★=84.1% (best)



(e) DermaMNIST Convergence Map
 $\square$ =trivial $\square$ =real learning



(f) PathMNIST AUNRC Ranking
 (standard noise-robustness metric)

